

CS 170

DISCUSSION 4

GRAPHS AND PATHS

Raymond Chan
UC Berkeley Fall 17

DEPTH FIRST SEARCH

- Explore all and only nodes reachable from current node.

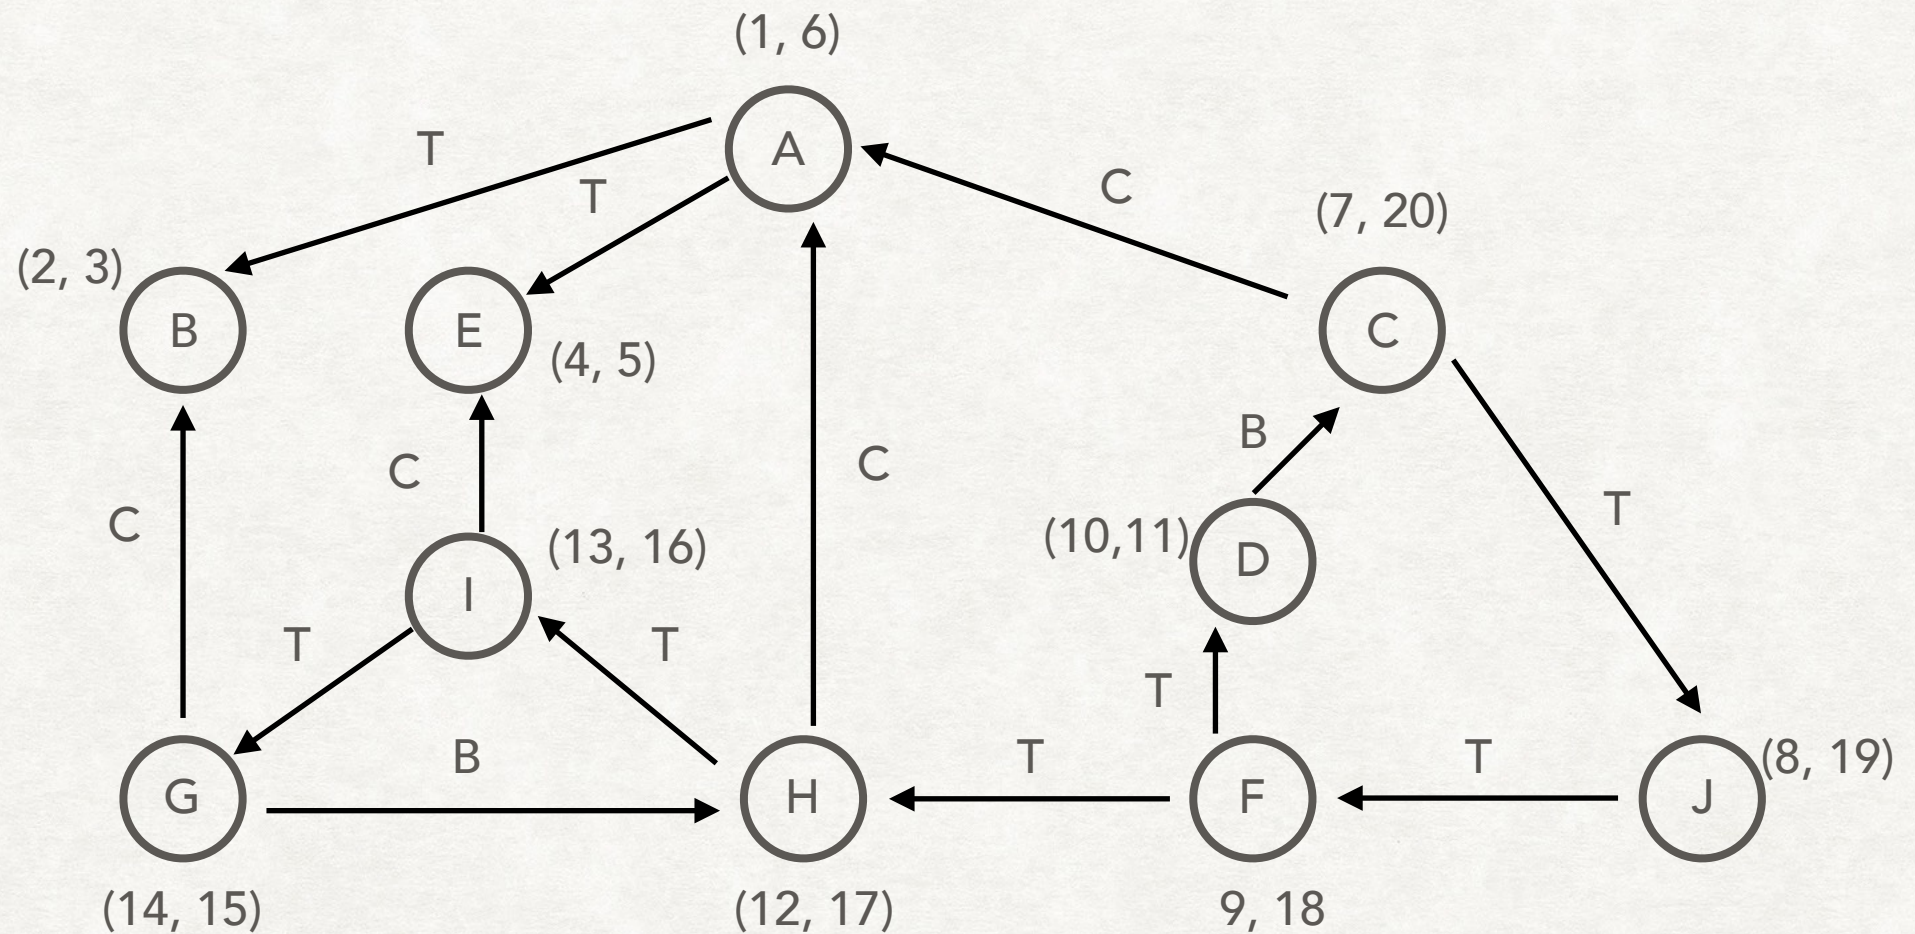
```
recursive_DFS (G, v):  
    previsit(v)  
    mark v as visited  
    for all v's neighbors w:  
        if vertex w has not been visited:  
            recursive_DFS (G, w)  
    postvisit(v)
```

```
iterative_DFS (G, v):  
    stack.push(v)  
    while stack not empty:  
        v = pop from stack  
        if v not visited  
            mark v as visited:  
            for all v's neighbors w:  
                push vertex w to stack
```

- Visits all vertices once. Uses all edges once.
- $O(|V| + |E|)$

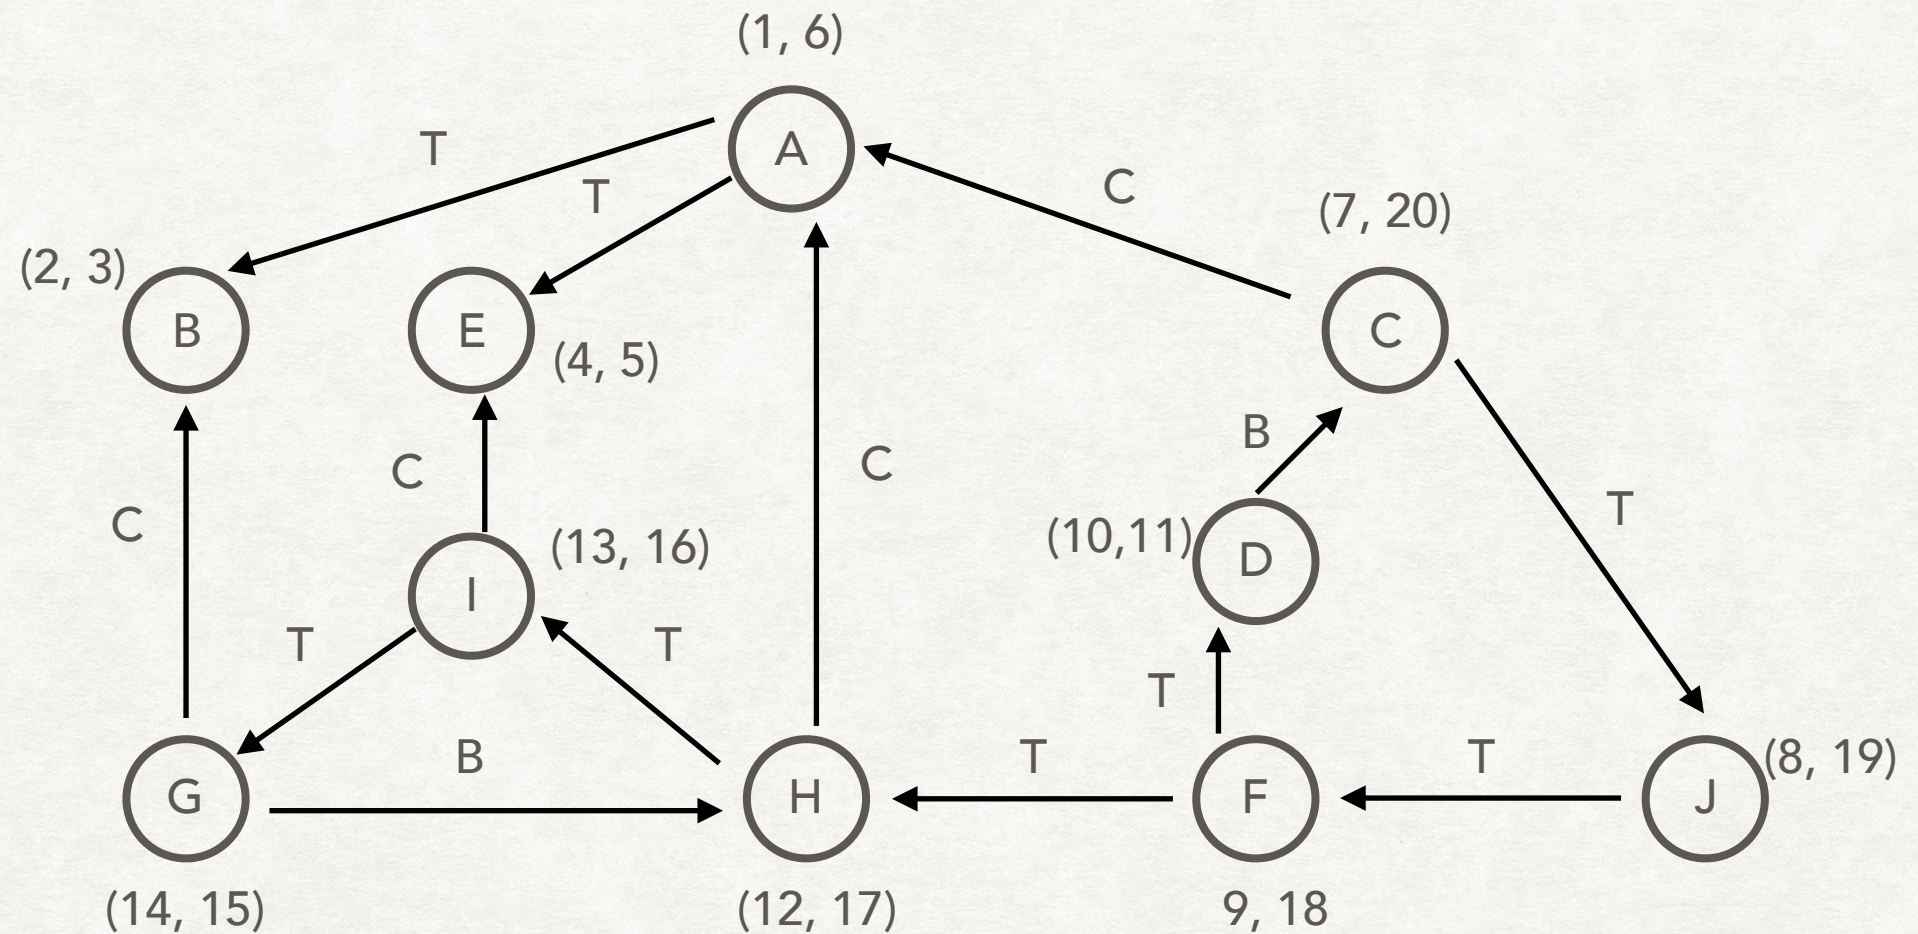
DEPTH FIRST SEARCH

- From last discussion



PREVISIT & POSTVISIT

- Previsit() and Postvisit() increments a global count



PREVISIT & POSTVISIT T/F

- If (u, v) is an edge in an indirect graph and during DFS, $\text{post}(v) < \text{post}(u)$, then u is an ancestor of v in the DFS tree.

PREVISIT & POSTVISIT T/F

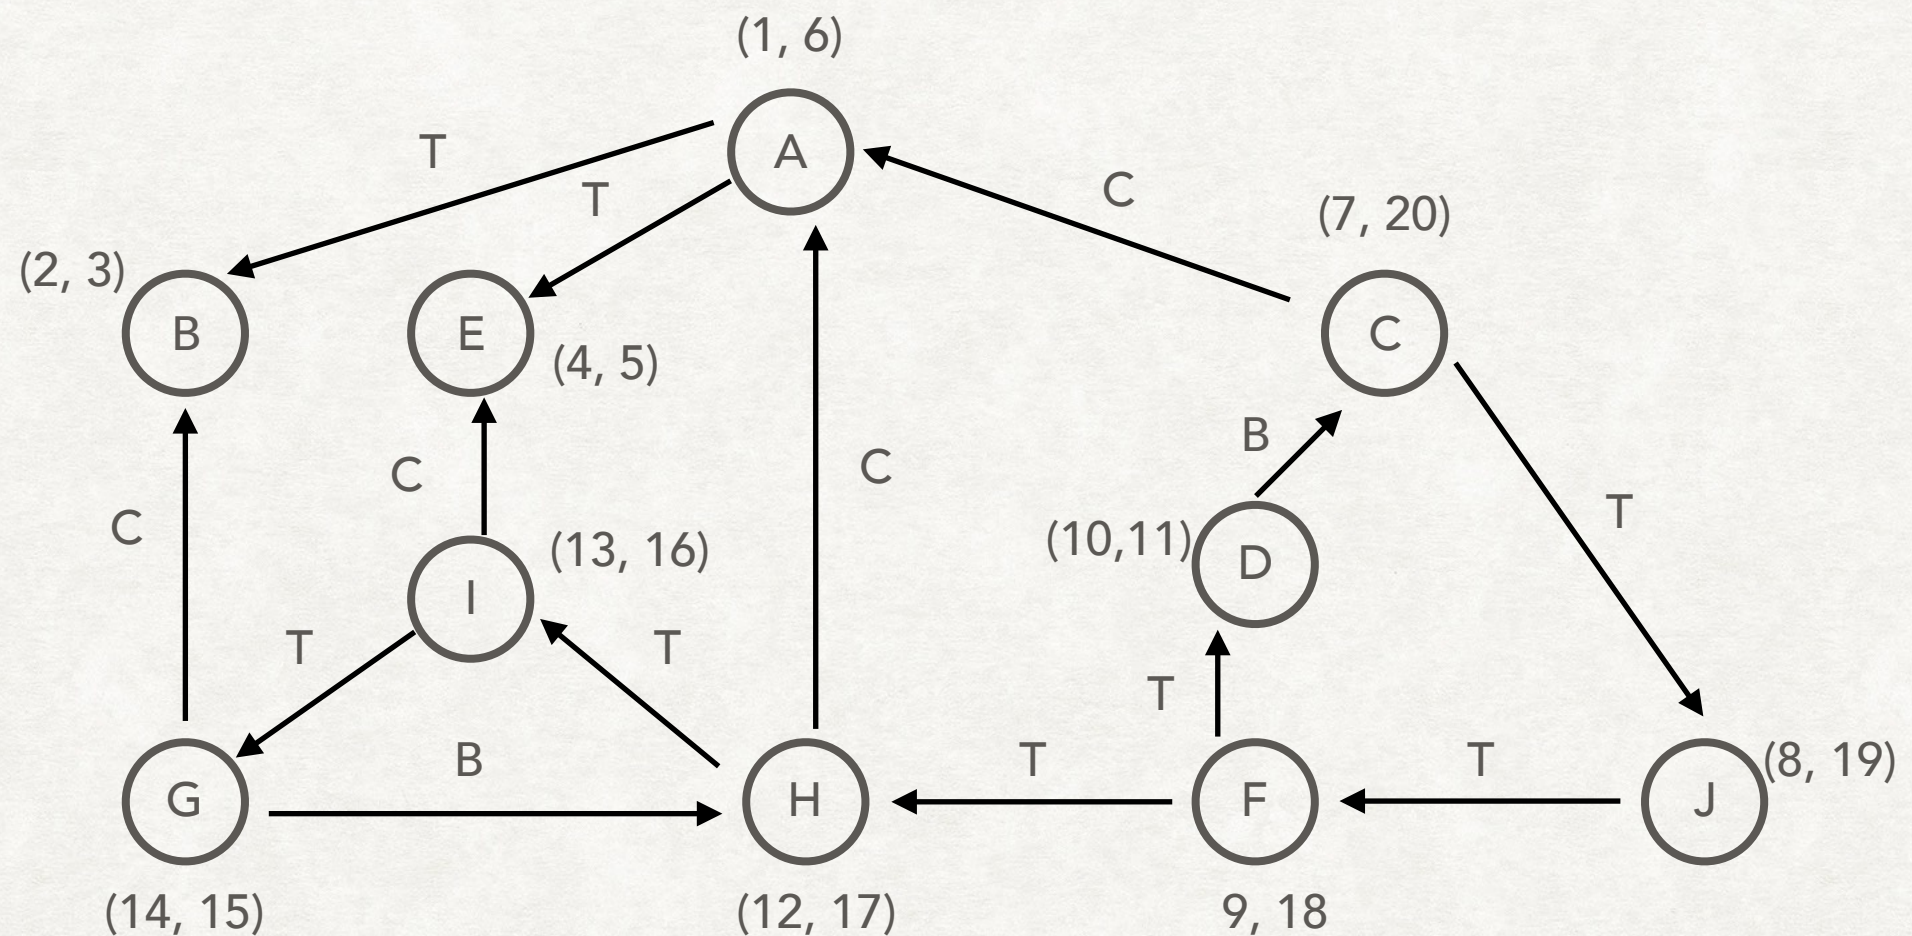
- If (u, v) is an edge in an indirect graph and during DFS, $\text{post}(v) < \text{post}(u)$, then u is an ancestor of v in the DFS tree.
- 2 cases (since $\text{pre} < \text{post}$):
 - $\text{pre}(u) < \text{pre}(v) < \text{post}(v) < \text{post}(u)$
 - u is an ancestor v . Explore u 's neighbors before postvisiting u

PREVISIT & POSTVISIT T/F

- If (u, v) is an edge in an indirect graph and during DFS, $\text{post}(v) < \text{post}(u)$, then u is an ancestor of v in the DFS tree.
- 2 cases (since $\text{pre} < \text{post}$):
 - $\text{pre}(u) < \text{pre}(v) < \text{post}(v) < \text{post}(u)$
 - u is an ancestor v . Explore u 's neighbors before postvisiting u
 - $\text{pre}(v) < \text{post}(v) < \text{pre}(u) < \text{post}(u)$
 - Looks at all of v 's neighbors before looking at u .
 - Contradiction since there is an edge.
- True

PREVISIT & POSTVISIT

- For any two nodes, u, v , $[\text{pre}(u), \text{post}(u)]$ and $[\text{pre}(v), \text{post}(v)]$ are either disjoint or one is contained in the other.

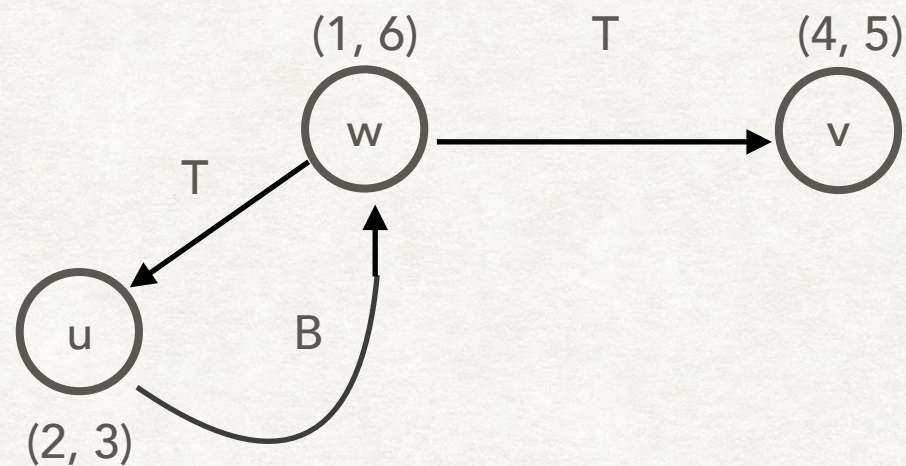


PREVISIT & POSTVISIT T/F

- In a directed graph, if there is a path from u to v and $\text{pre}(u) < \text{pre}(v)$ then u is an ancestor of v in the DFS tree.

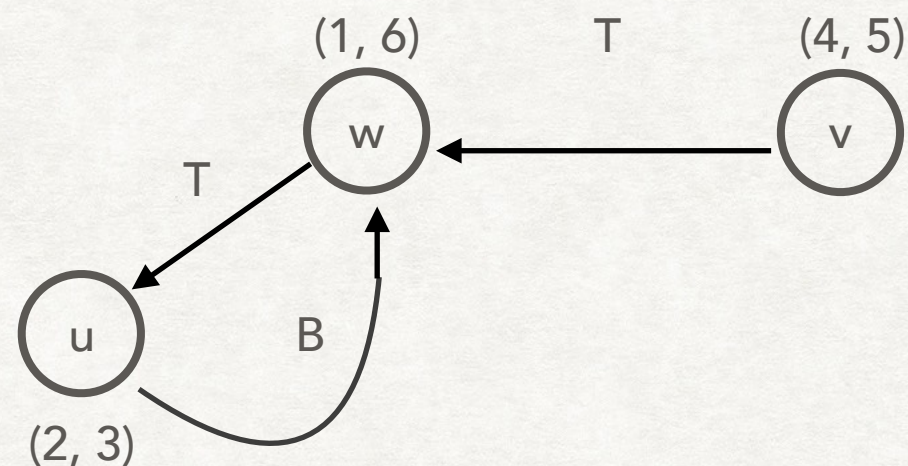
PREVISIT & POSTVISIT T/F

- In a directed graph, if there is a path from u to v and $\text{pre}(u) < \text{pre}(v)$ then u is an ancestor of v in the DFS tree.
- Consider the case when u and v are a common ancestor's direct children.



PREVISIT & POSTVISIT T/F

- In a directed graph, if there is a path from u to v and $\text{pre}(u) < \text{pre}(v)$ then u is an ancestor of v in the DFS tree.
- Consider the case when u and v are a common ancestor's direct children.



- u gets visited first via w , who then visits v .
- Need information about post visit number
- False

PREVISIT & POSTVISIT T/F

- In any connected undirected graph G , there is a vertex whose removal leaves G connected.

PREVISIT & POSTVISIT T/F

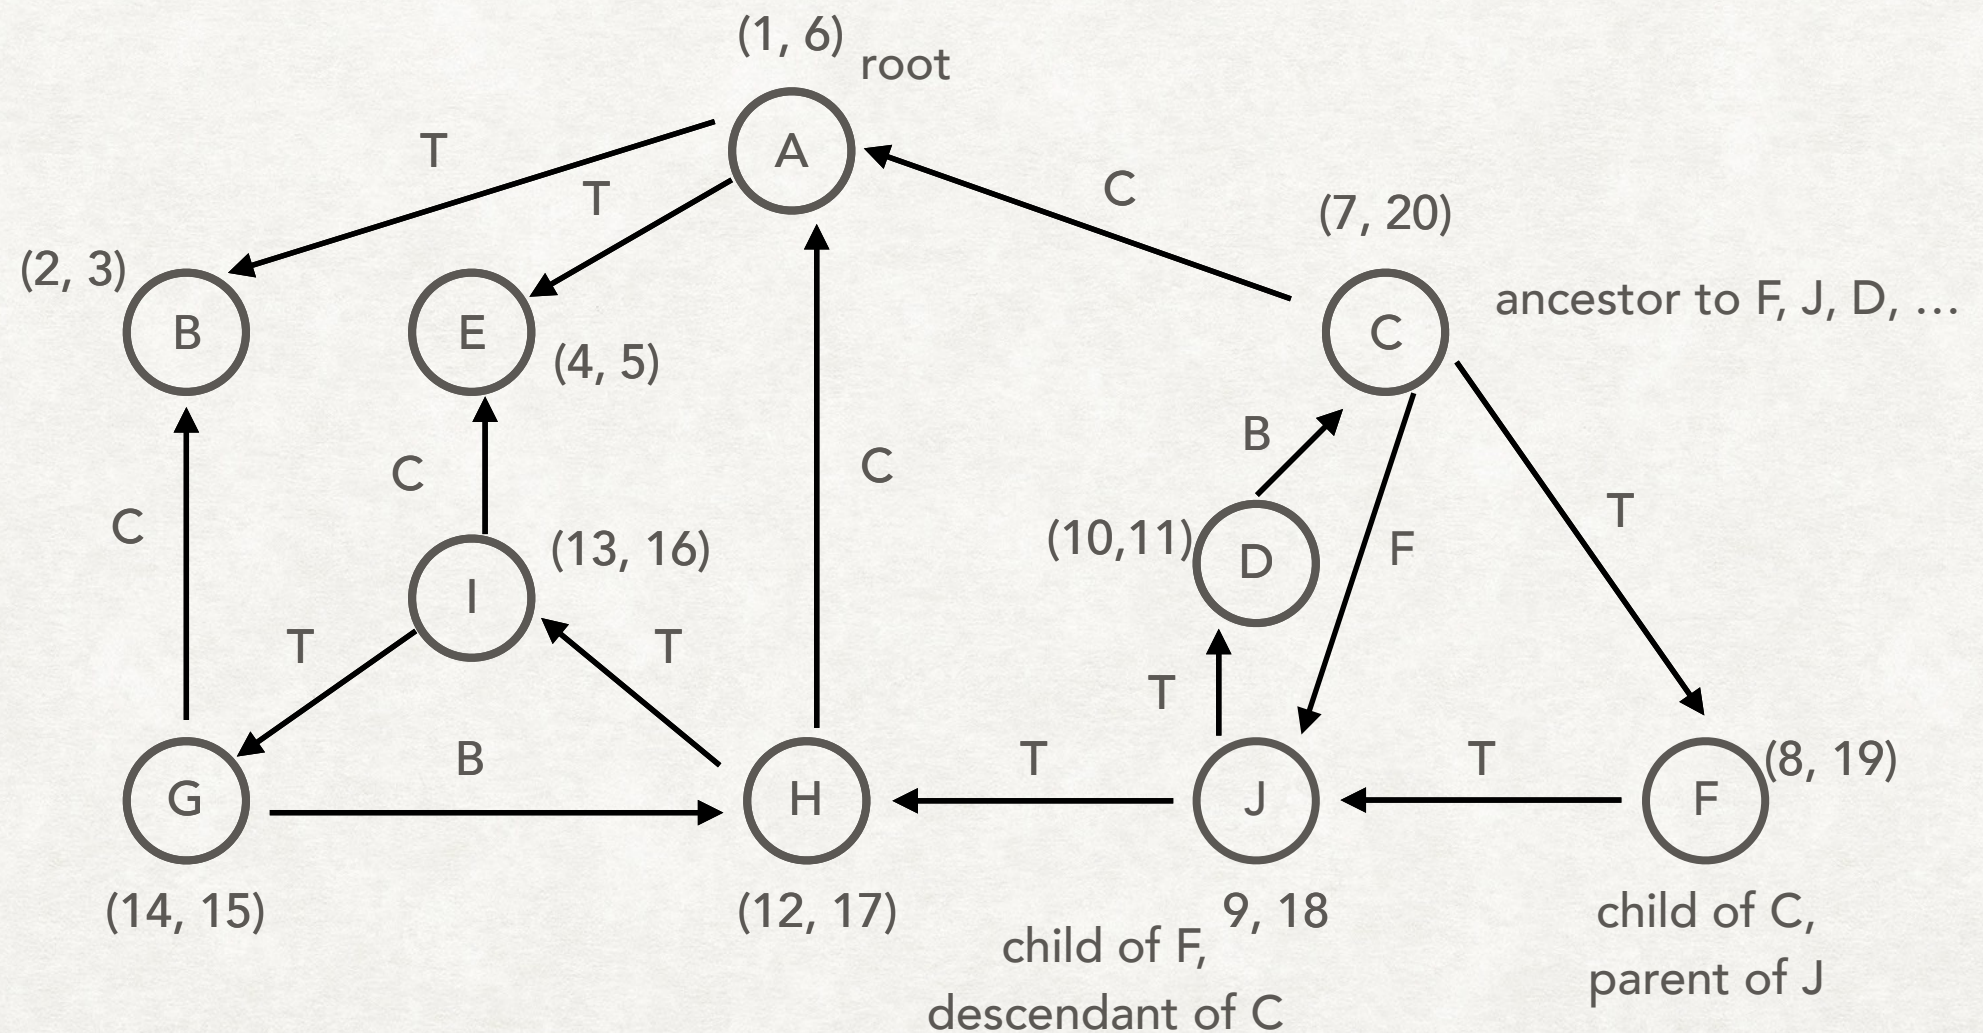
- In any connected undirected graph G , there is a vertex whose removal leaves G connected.
- True

PREVISIT & POSTVISIT T/F

- In any connected indirect graph G , there is a vertex whose removal leaves G connected.
- True
- Removing any leaf from a DFS tree of the graph.
- These leaves have only one neighbor, otherwise they won't be leaves.
- Neighbor is connected to at least one other node.

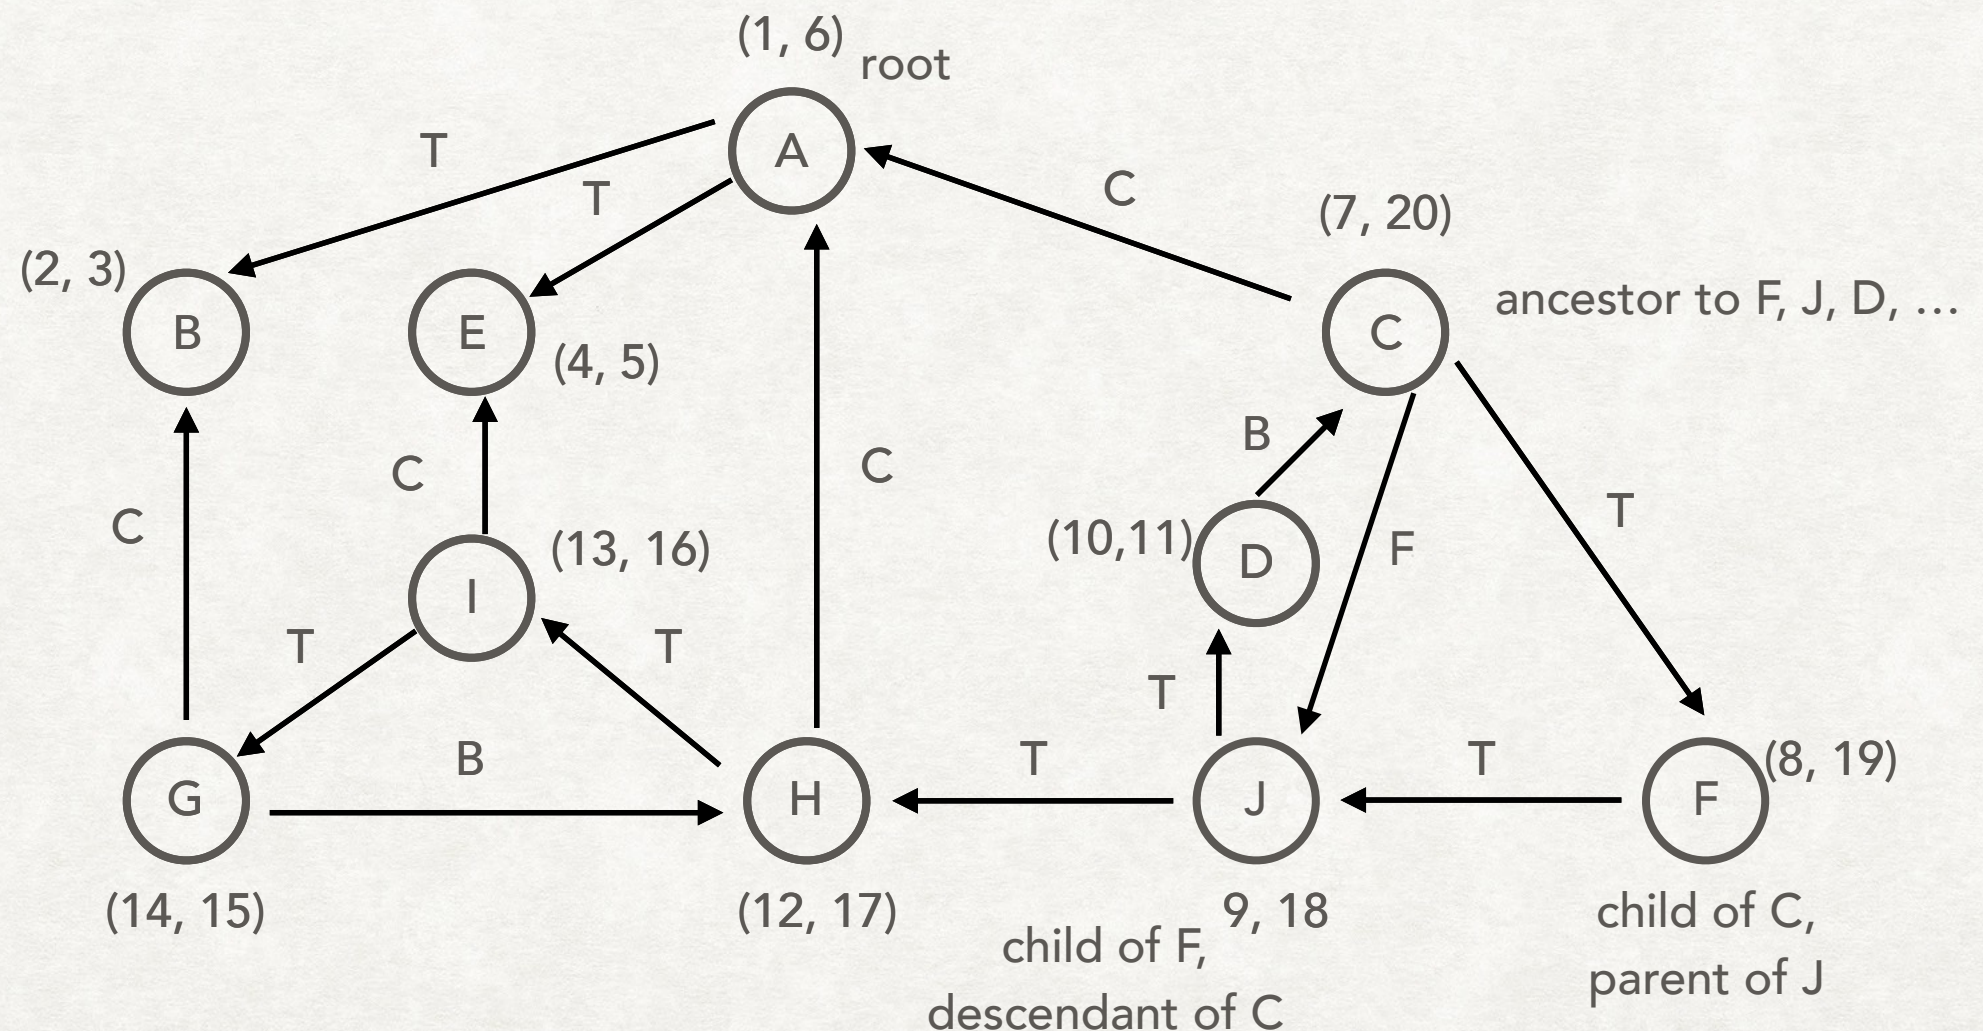
DEPTH FIRST SEARCH

- Tree edge (u, v) : part of DFS
- $\text{pre}(u) < \text{pre}(v) < \text{post}(v) < \text{post}(u)$



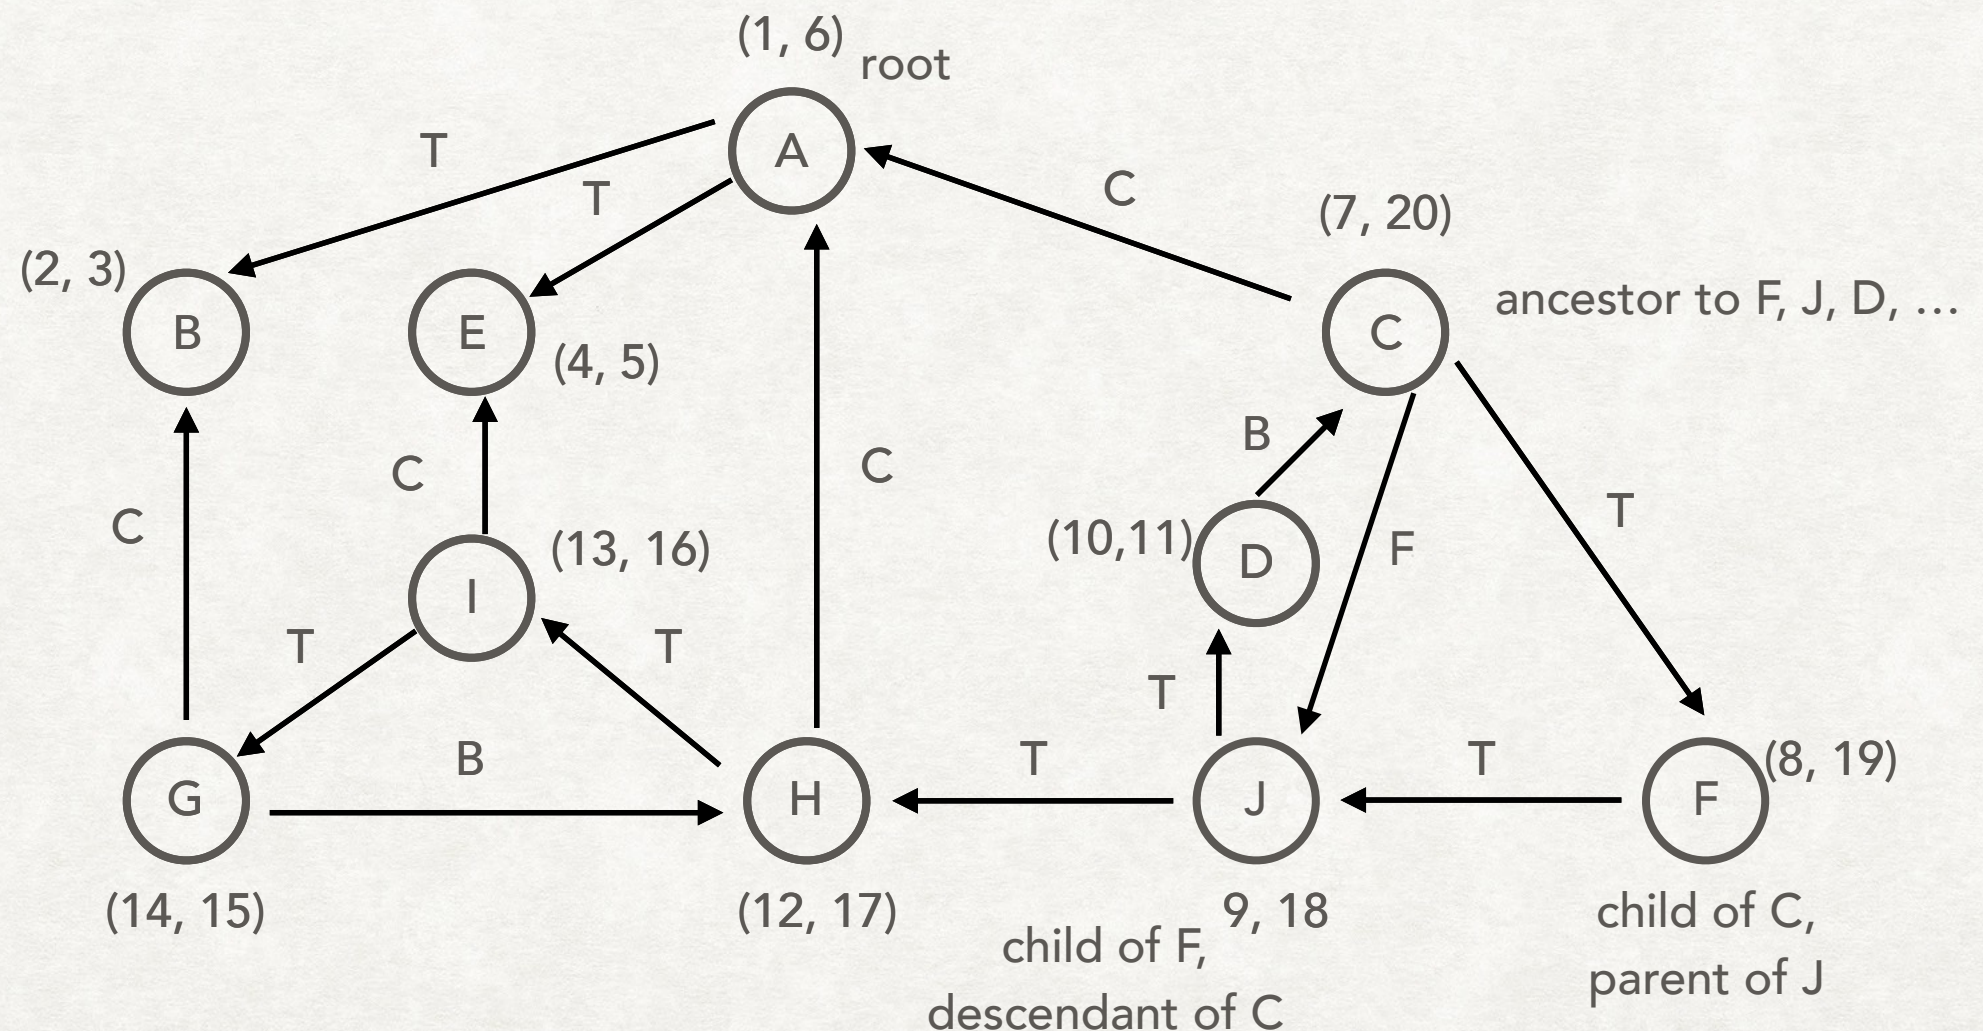
DEPTH FIRST SEARCH

- Forward edge (u, v) : leads non-child descendant
- $\text{pre}(u) < \text{pre}(v) < \text{post}(v) < \text{post}(u)$



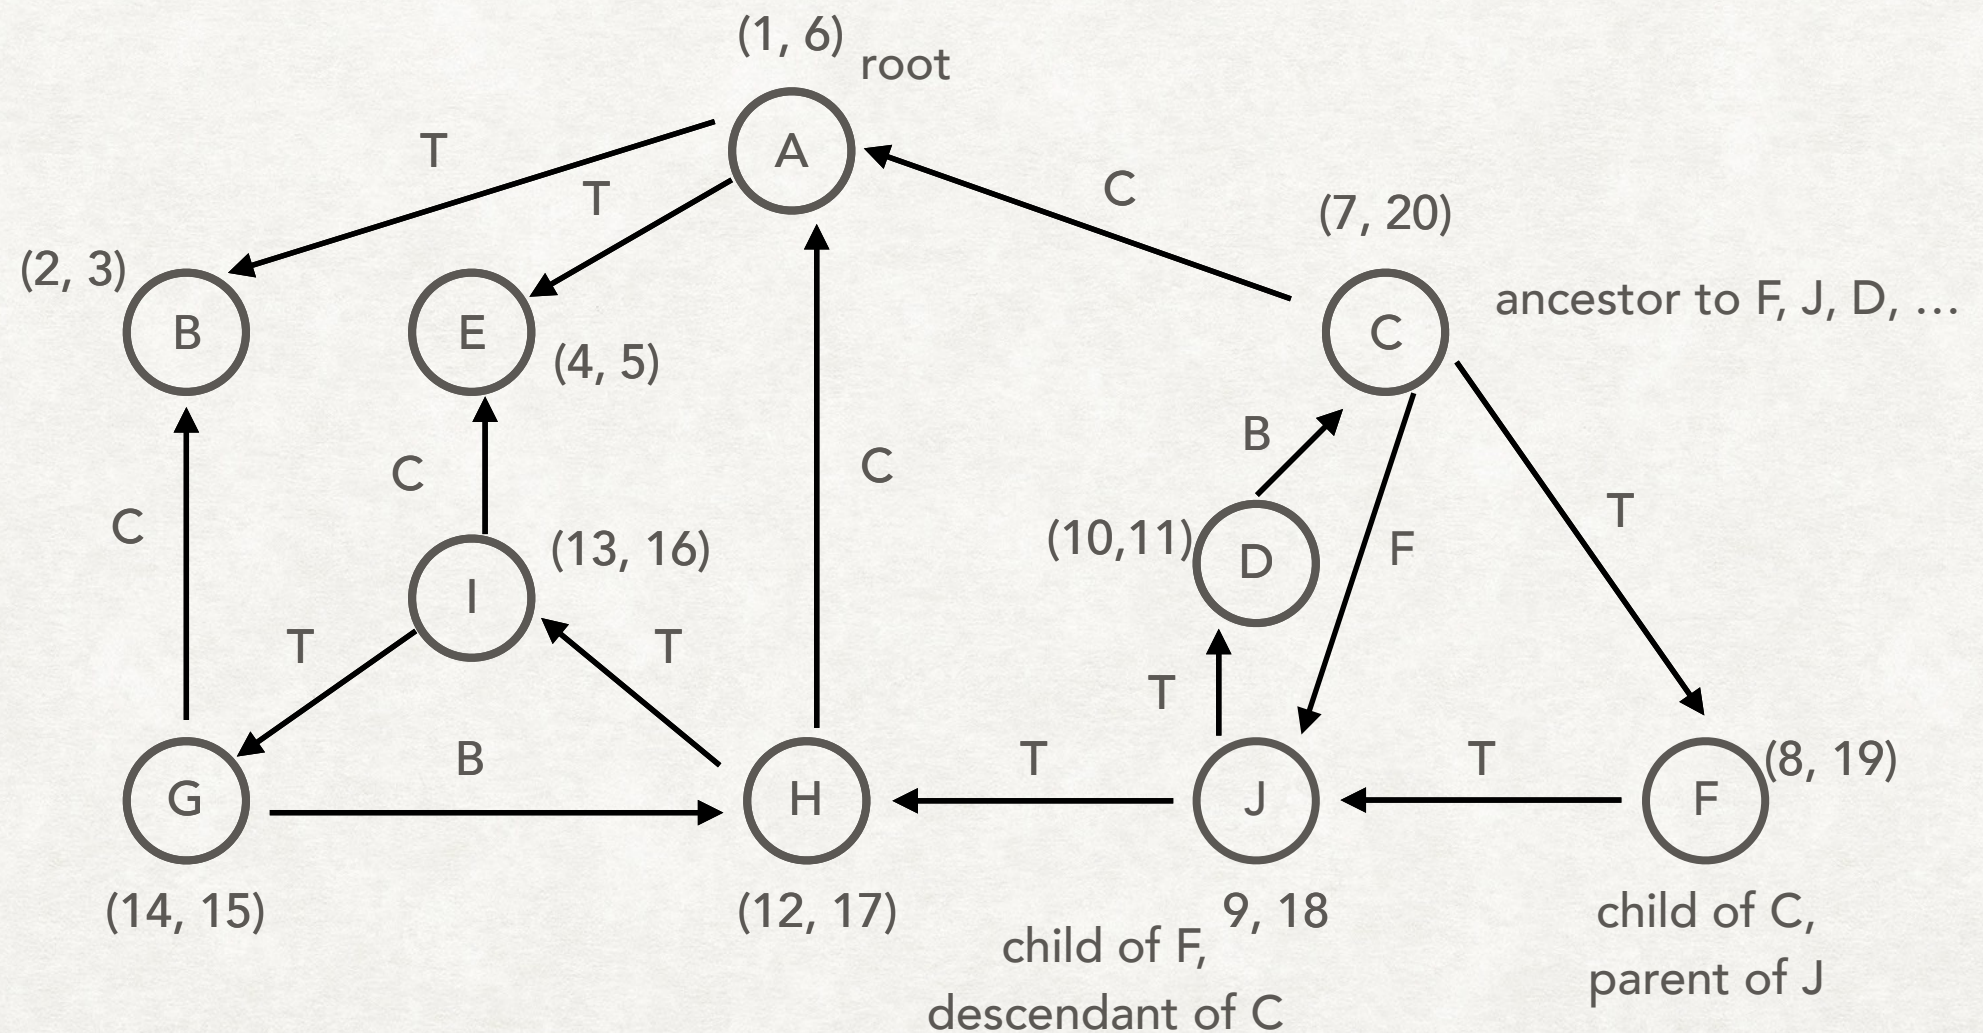
DEPTH FIRST SEARCH

- Cross edge (u, v) : leads neither descendant nor ancestor
- $\text{pre}(v) < \text{post}(v) < \text{pre}(u) < \text{post}(u)$



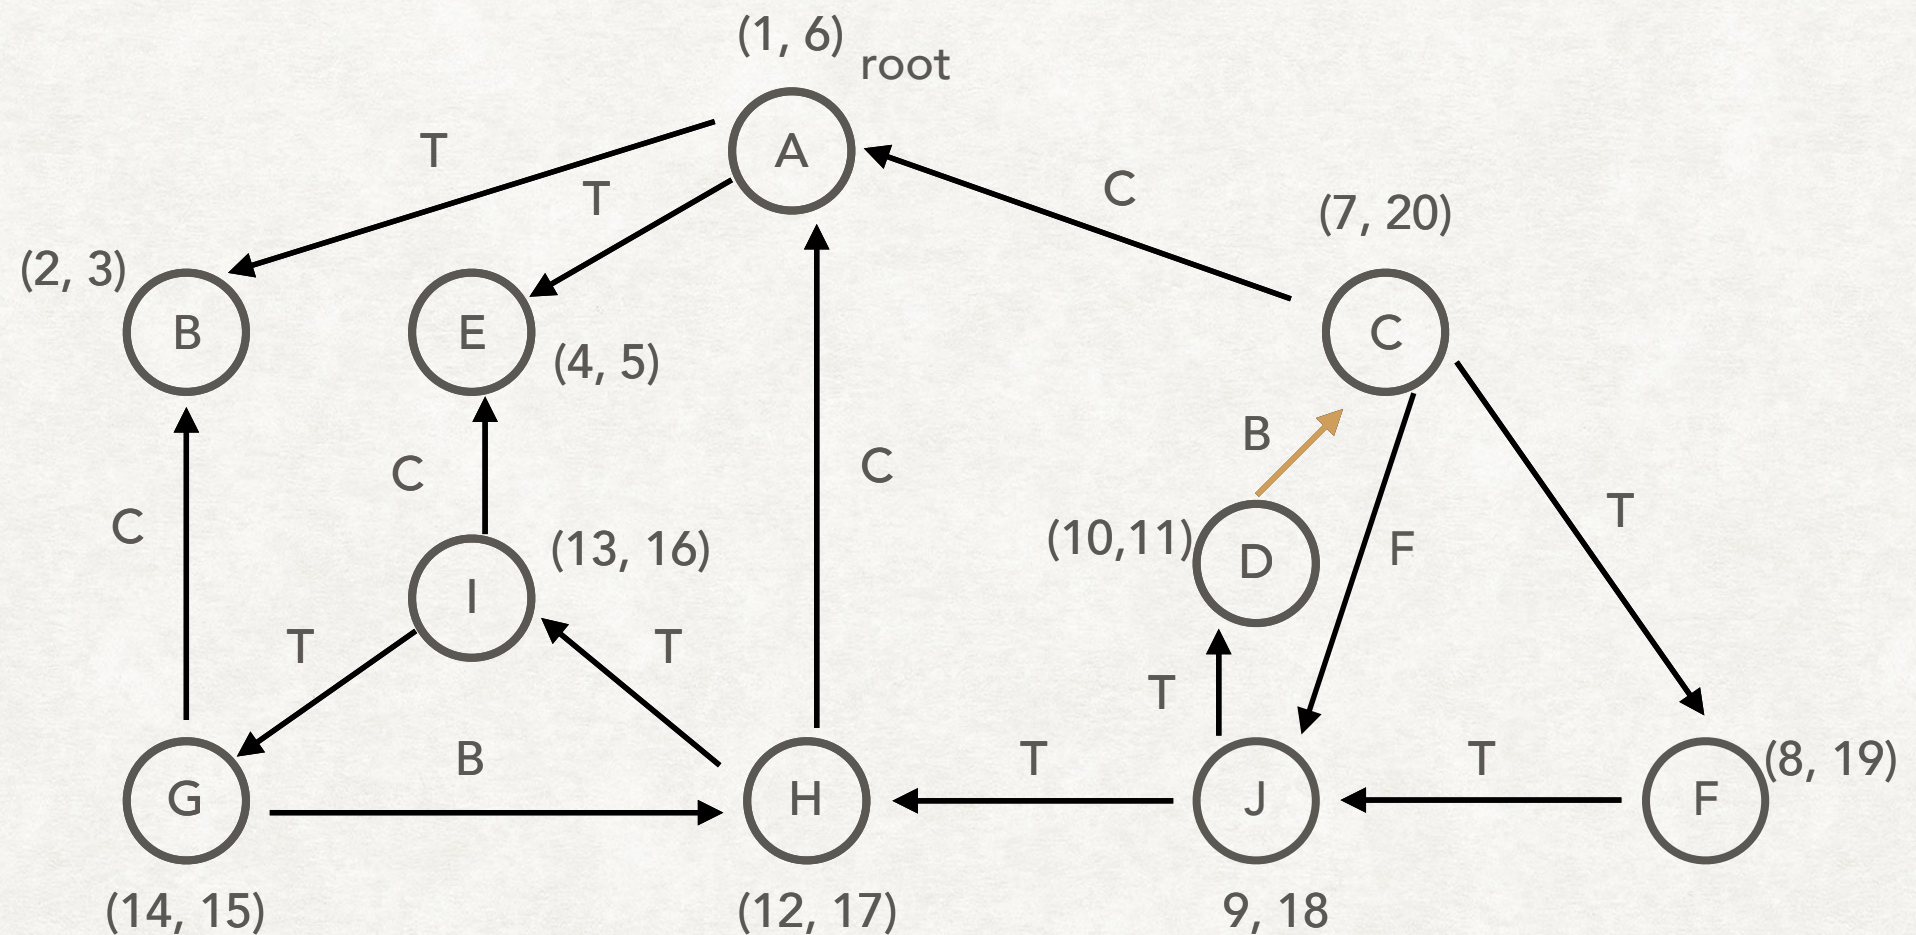
DEPTH FIRST SEARCH

- Back edge (u, v): leads to ancestor
- $\text{pre}(v) < \text{pre}(u) < \text{post}(u) < \text{post}(v)$



CYCLE DETECTION

- There is a cycle if and only if there is a back edge.
- Run DFS.

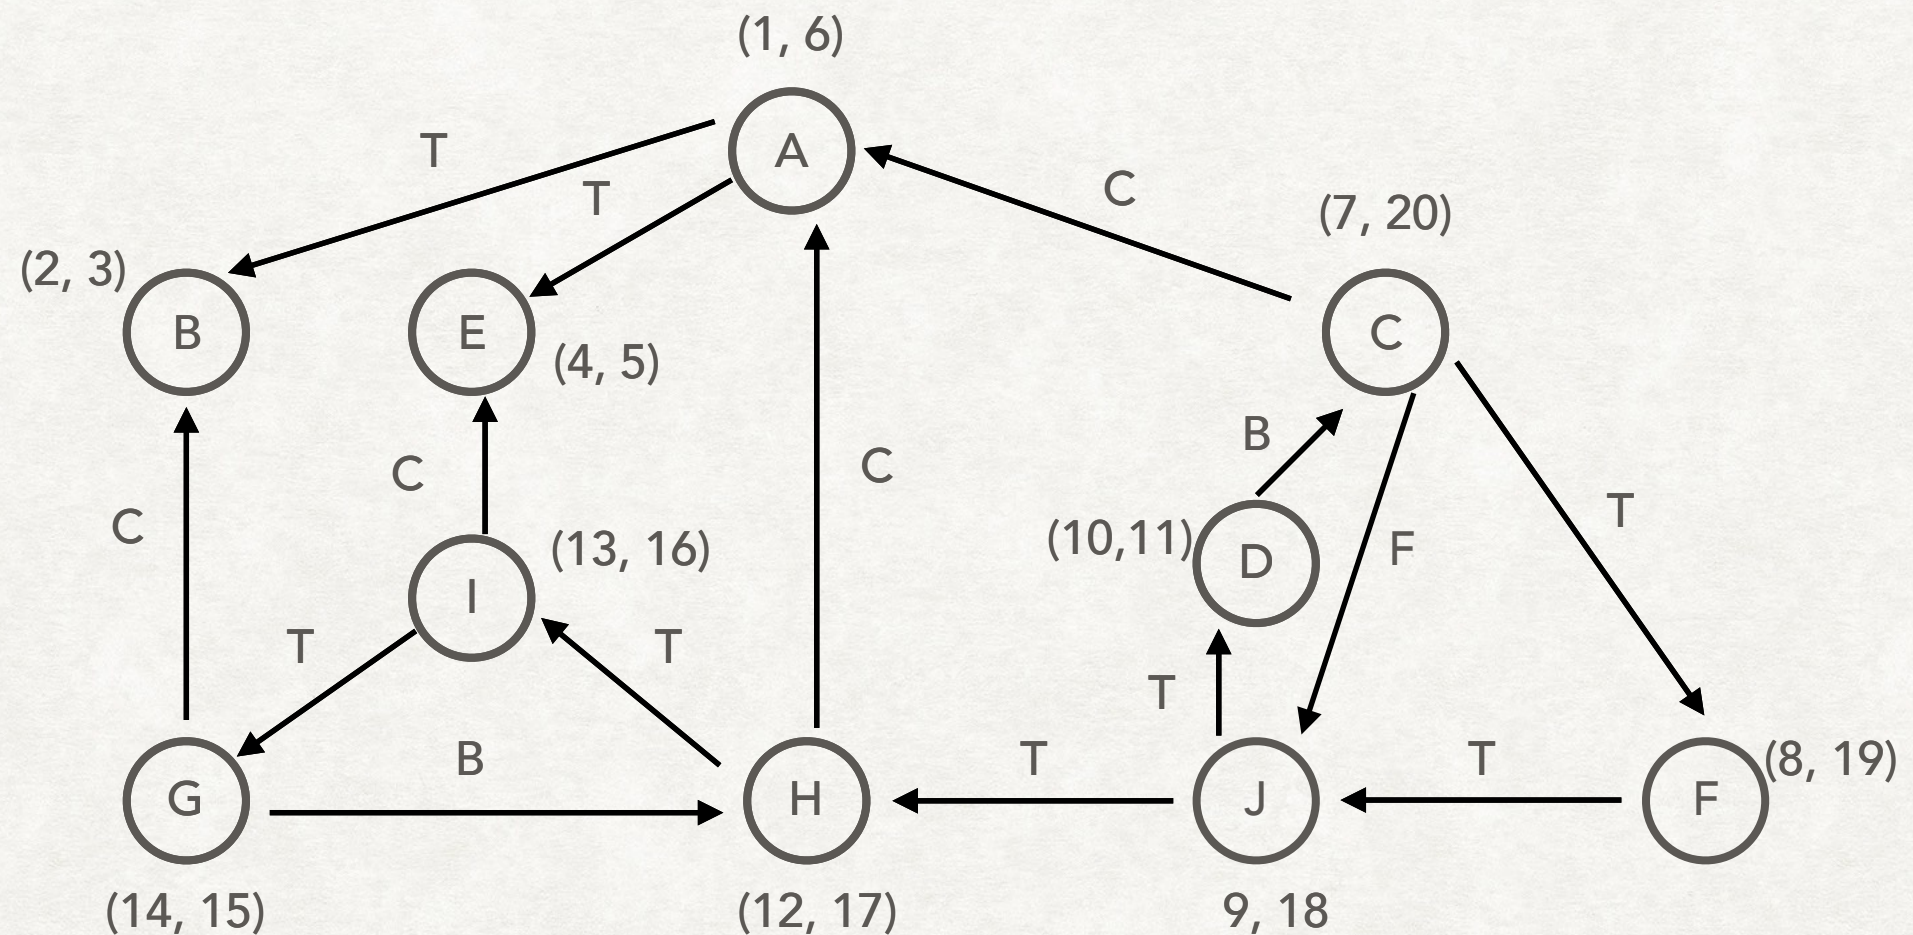


STRONGLY CONNECTED COMPONENTS

- Nodes u and v are strongly connected if there is a path from u to v and there is a path from v to u .
- Strongly connected components are a set of nodes that are strongly connected.
- Can visit every other node in the set.
- Only applies to directed graphs.
- Single case: single node.

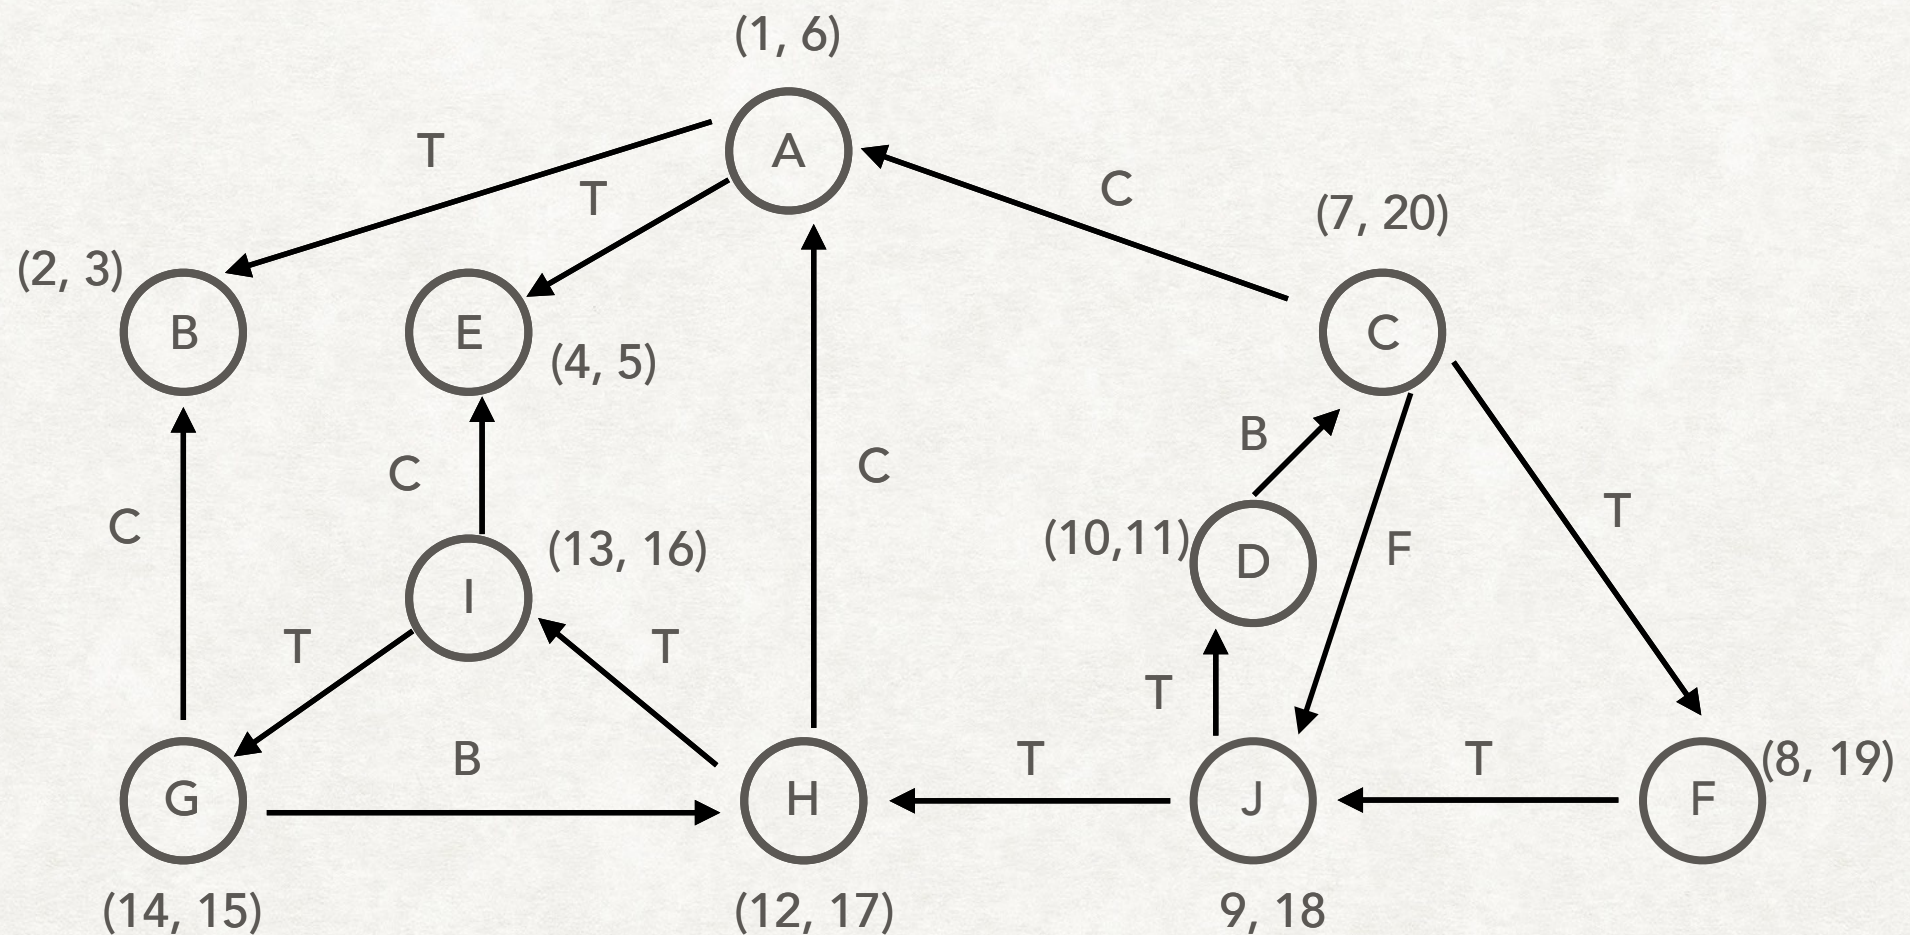
STRONGLY CONNECTED COMPONENTS

- Find the strongly connected components.



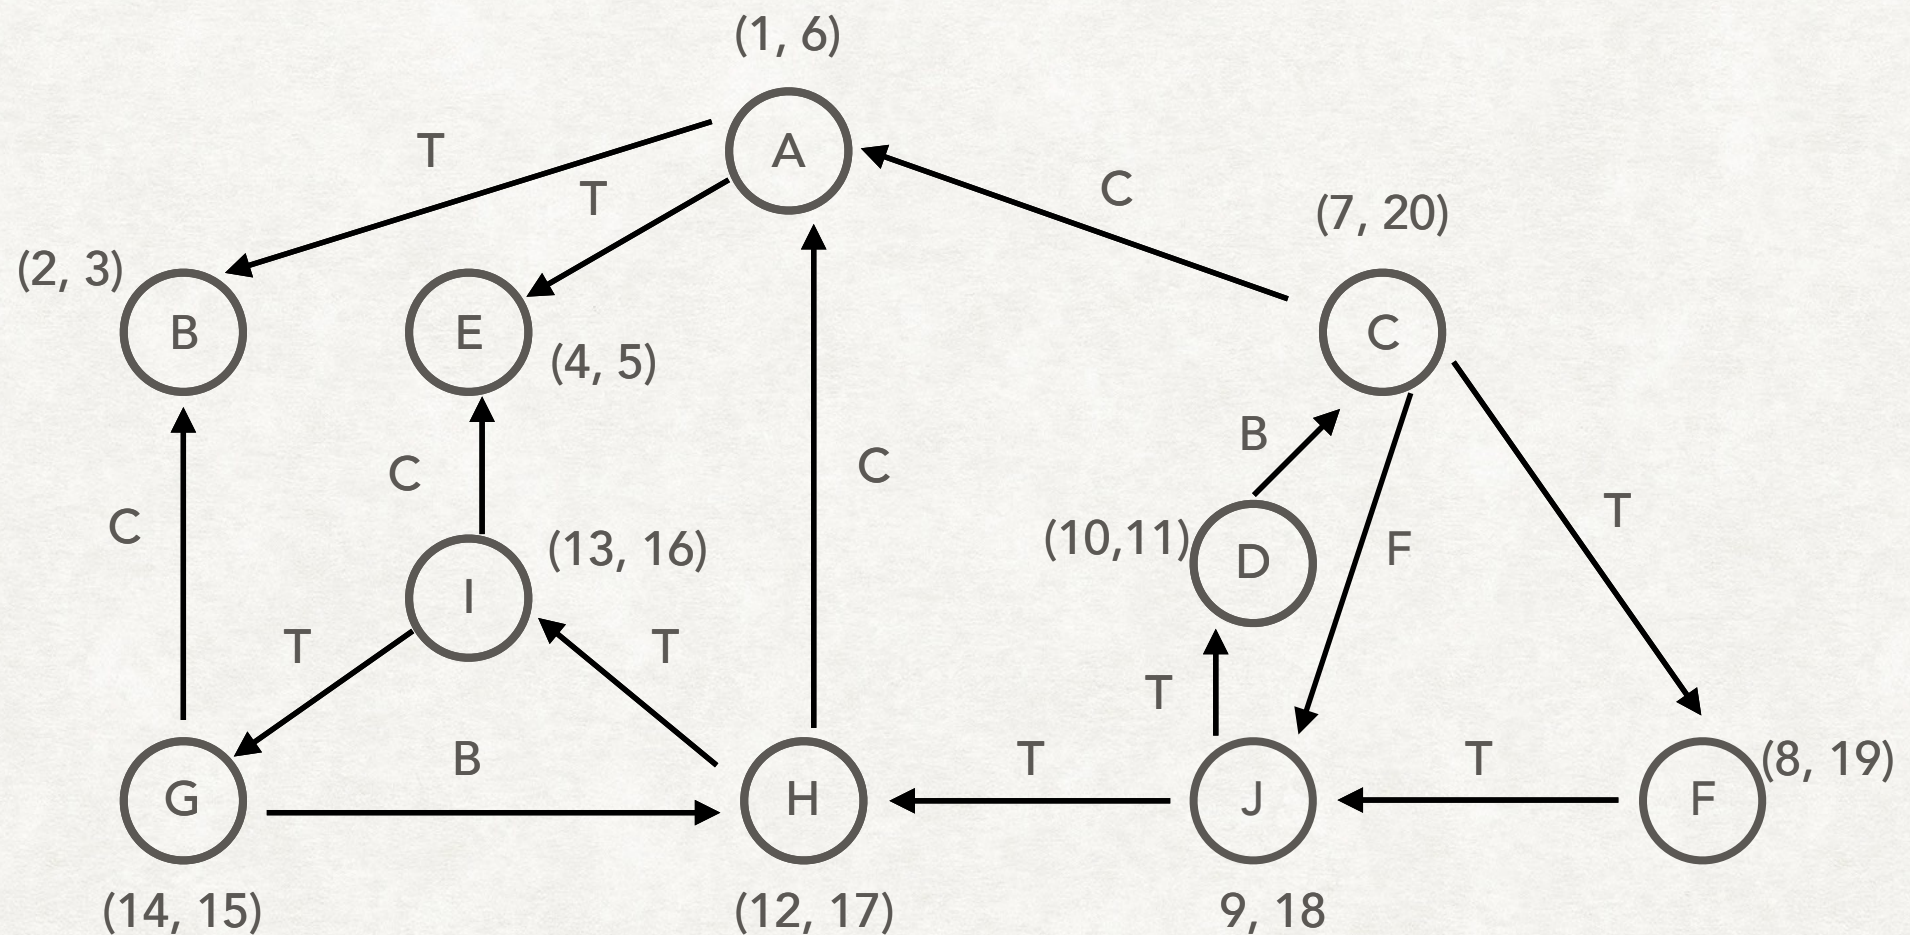
STRONGLY CONNECTED COMPONENTS

- Find the strongly connected components.
- $\{A\}$, $\{B\}$, $\{E\}$, $\{C, F, J, D\}$, $\{G, I, H\}$



STRONGLY CONNECTED COMPONENTS

- Visualize it
 - <https://www.cs.usfca.edu/~galles/JavascriptVisual/ConnectedComponent.html>



TOPOLOGICAL SORT

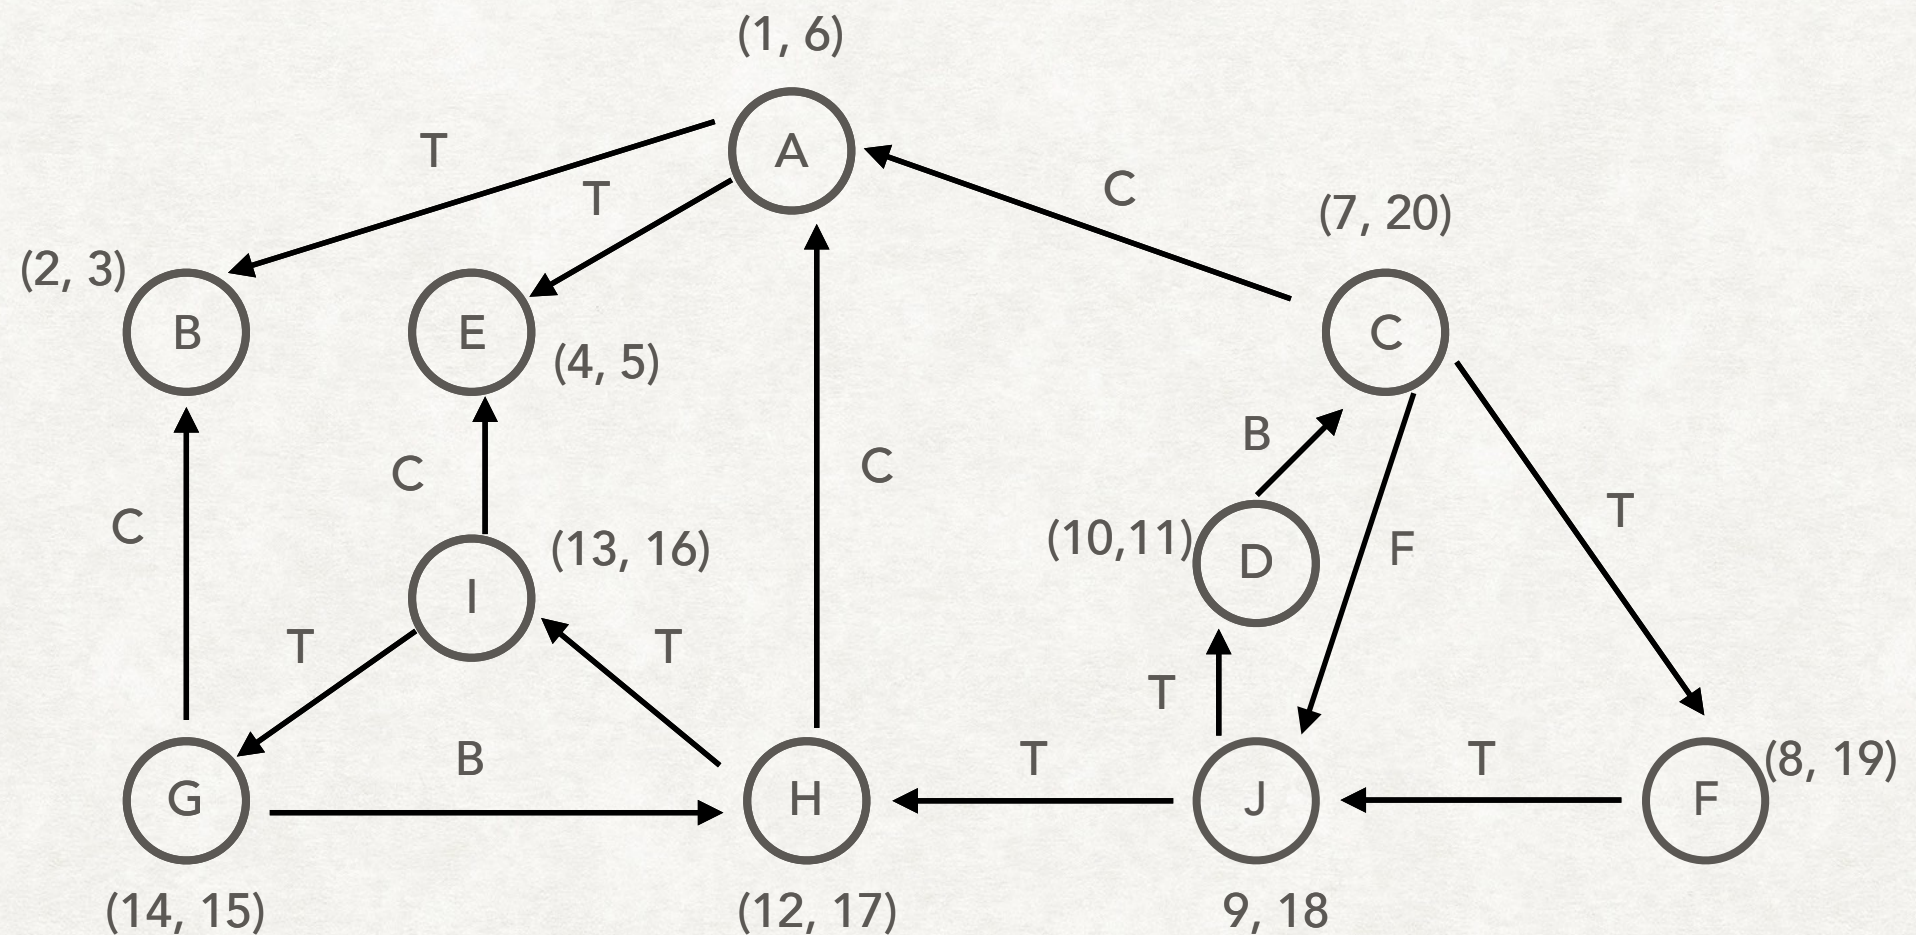
- Source and Sink in a Directed Acyclic Graph (DAG)
- Source:
 - Node has no incoming edges
 - Highest post order number
 - First nodes in topological ordering
- Sink
 - No outgoing edges
 - Lowest post order number
 - Last nodes in topological ordering
- Every DAG has at least one source and one sink

TOPOLOGICAL SORT

- Ordering of a directed graph's nodes v_1, v_2, \dots, v_n such that for every edge (v_i, v_j) , $i < j$.
- Edge arrows go one direction.
- Application: scheduling jobs if order is required.

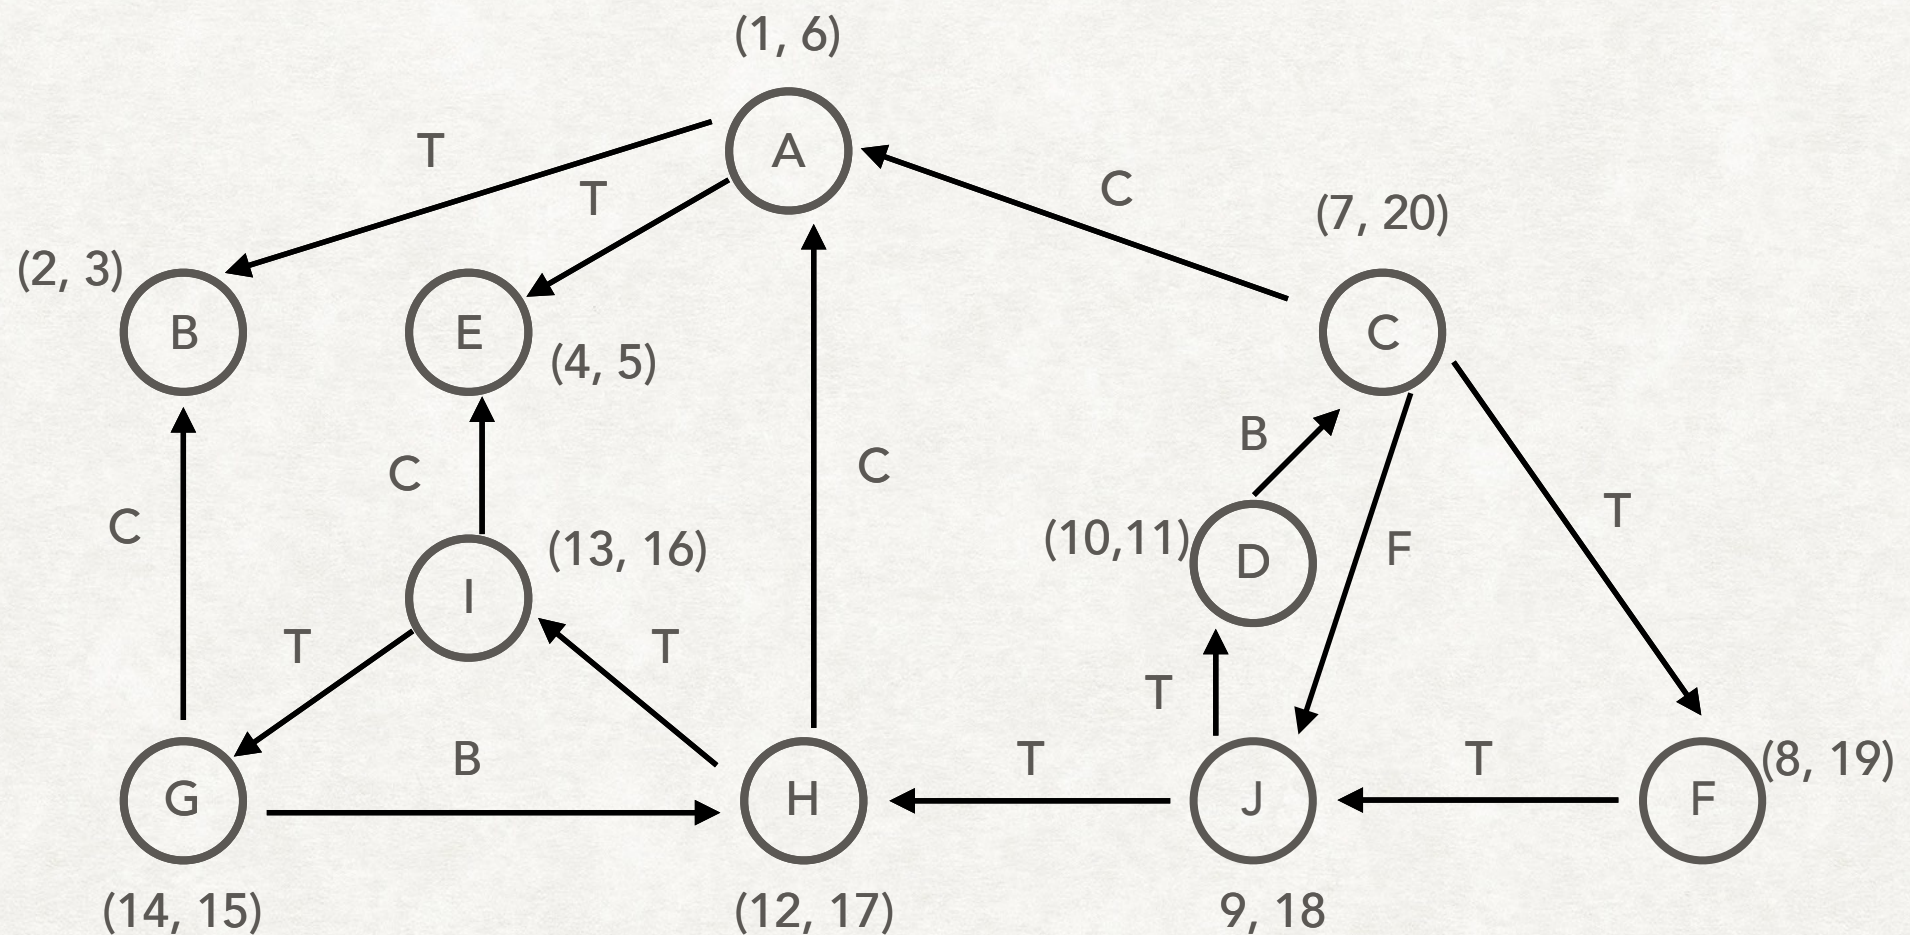
TOPOLOGICAL SORT

- Sort the Strongly Connected Components into a DAG



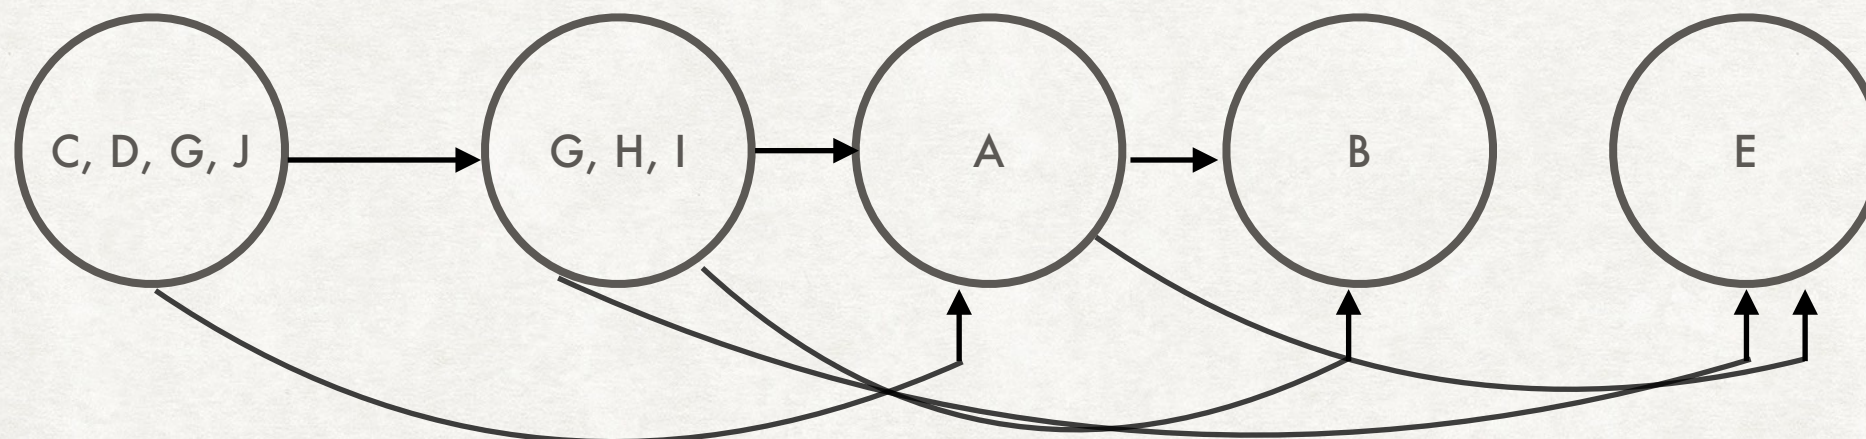
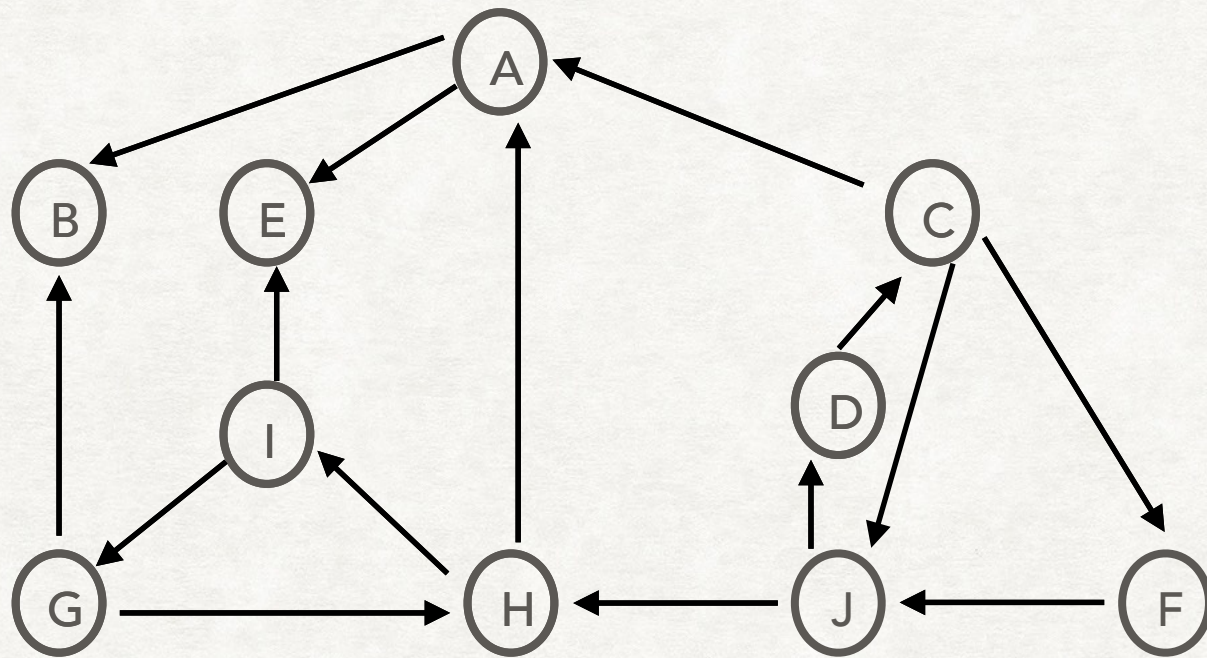
TOPOLOGICAL SORT

- Sort the Strongly Connected Components into a DAG
- {A}, {B}, {E}, {C, F, J, D}, {G, I, H}



TOPOLOGICAL SORT

- {A}, {B}, {E}, {C, F, J, D}, {G, I, H}



ALGORITHMS

SCC(G):

Reverse edges of graph $\rightarrow G^R$

DFS on G^R

Run DFS on G in reverse post
order number from G^R

- Reverse graph prevents crossing SCCs
- Tarjan's SCC algorithm

Linearize (G, v):

push all sources to v

while stack not empty:

v = pop from stack

for all v's neighbors w:

remove edge (v, w)

if w becomes a sources:

push vertex w to stack

Add v to result

Sort by reverse post order number

BREADTH FIRST SEARCH

- Explore node u 's neighbors, then all vertices that are adjacent to u 's neighbors, and so on.
- Can keep distance values to find how many edges a node is away from the root.
- Visits all vertices once.
- Uses all edges once.
- $O(|V| + |E|)$

```
iterative_BFF (G, v):  
    d[v] = 0  
    queue.enqueue(v)  
    while queue not empty:  
        v = pop from queue  
        for all v's neighbors w:  
            if w not visited  
                mark w as visited:  
                d[w] = d[v] + 1  
                queue.enqueue(w)
```


DIJKSTRA'S SHORTEST PATH

- Find shortest path from s to all other vertices.
- Once we have computed the shortest path for a vertex, we don't revisit it again.

```
dijkstra (G, s):  
    d[v] = infinity  
    d[s] = 0  
    prev[s] = s  
    PQ.add(G.V, infinity)  
    PQ.add(s, 0)  
    while PQ not empty:  
        u = PQ.DeleteMin()  
        for edge (u, v):  
            if d[v] > d[u] + w(u, v):  
                d[v] = d[u] + w[u, v]  
                prev[v] = u  
                PQ.DecreaseKey(v, d[v])
```


DIJKSTRA'S SHORTEST PATH

- Keep fringe vertices.
- Look at neighbors and update distance if there is a better path.
- When we pop off a node from Priority Queue, the distance is the shortest path from s to this node so far.

```
dijkstra (G, s):  
    d[v] = infinity  
    d[s] = 0  
    prev[s] = s  
    PQ.add(G.V, infinity)  
    PQ.add(s, 0)  
    while PQ not empty:  
        u = PQ.DeleteMin()  
        for edge (u, v):  
            if d[v] > d[u] + w(u, v):  
                d[v] = d[u] + w[u, v]  
                prev[v] = u  
                PQ.DecreaseKey(v, d[v])
```


DIJKSTRA'S SHORTEST PATH

Demo

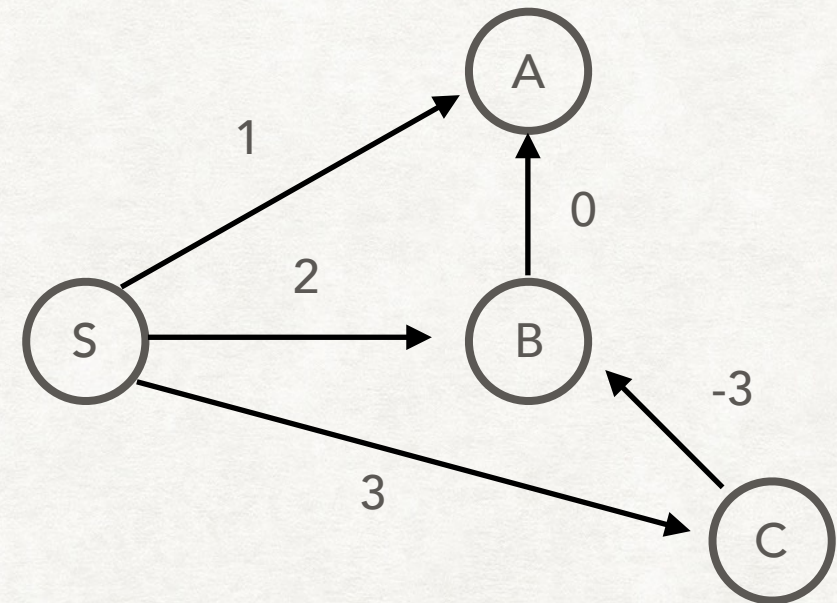
```
dijkstra (G, s):  
    d[v] = infinity  
    d[s] = 0  
    prev[s] = s  
    PQ.add(G.V, infinity)  
    PQ.add(s, 0)  
    while PQ not empty:  
        u = PQ.DeleteMin()  
        for edge (u, v):  
            if d[v] > d[u] + w(u, v):  
                d[v] = d[u] + w[u, v]  
                prev[v] = u  
                PQ.DecreaseKey(v, d[v])
```

$O(|V|) * \text{DeleteMin} + O(|E|) * \text{DecreaseKey}$

Typically use binary heap for priority queue

$O(|V| + |E|) * \log(|V|)$

NEGATIVE EDGE WEIGHTS



```
dijkstra (G, s):  
    d[v] = infinity  
    d[s] = 0  
    prev[s] = s  
    PQ.add(G.V, infinity)  
    PQ.add(s, 0)  
    while PQ not empty:  
        u = PQ.DeleteMin()  
        for edge (u, v):  
            if d[v] > d[u] + w(u, v):  
                d[v] = d[u] + w[u, v]  
                prev[v] = u  
                PQ.DecreaseKey(v, d[v])
```


NEGATIVE EDGE WEIGHTS

Start at S

Process A. $d[a] = 1$

Process B. $d[b] = 2$

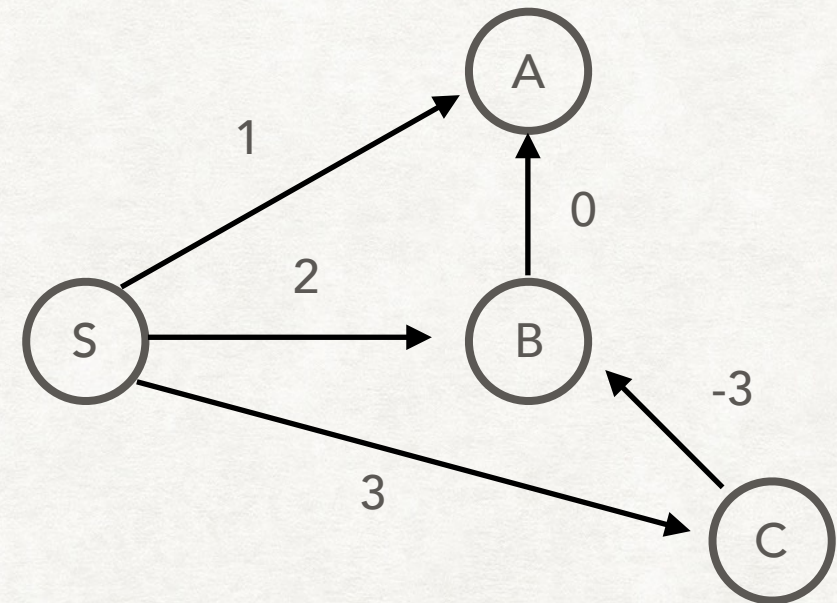
Process C. $d[c] = 3$

Process A. $d[a] = 1$.

Process B. $d[b] = 2$. $d[a] = 1$

Process C. $d[c] = 3$. $d[b] = 0$

New distance for B, but B not in PQ



```
dijkstra (G, s):  
  d[v] = infinity  
  d[s] = 0  
  prev[s] = s  
  PQ.add(G.V, infinity)  
  PQ.add(s, 0)  
  while PQ not empty:  
    u = PQ.DeleteMin()  
    for edge (u, v):  
      if d[v] > d[u] + w(u, v):  
        d[v] = d[u] + w[u, v]  
        prev[v] = u  
        PQ.DecreaseKey(v, d[v])
```